

**墨墨导读：**在Oracle 11g 版本中可能出现由于JDBC bug导致sql绑定变量无法共享，过期游标过多的情况，此时如果发生大量并发业务，很有可能造成异常library cache lock等待事件，造成数据库突发性能问题。在此，我们分享一次由jdbc bug和绑定变量长度问题共同“作案”引发数据库性能故障的案例。

本文节选自《云和恩墨技术通讯》（12月刊）

下载链接：<https://www.modb.pro/doc/1593>



library cache lock等待事件是Oracle数据库较为常见的等待事件之一，在之前的几次月刊中，我们也提到过产生library cache lock等待出现的原因有很多，如登录密码错误尝试过多、热表收集统计信息和SQL解析失败等。

在Oracle 11g 版本中可能出现由于JDBC bug导致sql绑定变量无法共享，过期游标过多的情况，此时如果发生大量并发业务，很有可能造成异常library cache lock等待事件，造成数据库突发性能问题。在此，我们分享一次由jdbc bug和绑定变量长度问题共同“作案”引发数据库性能故障的案例，供各位参考。

## 问题描述

---

2019年10月11号晚22:00分左右，运维人员对生产系统数据库进行清理历史分区操作，执行近100个分区删除操作后（22:05左右）发现该数据库压力飙升，维护人员紧急停止历史分区清理操作，发现大量业务数据插入（INSERT）缓慢。

查看故障期间数据库发现大量library cache lock等待，数据库活动会话飙升至1000以上，数据库响应非常缓慢，业务受到严重影响。

## 问题分析

---

从故障期间ASH的整体运行情况看：

7	2019-10-11 21:59:37	3
8	2019-10-11 22:00:08	39
9	2019-10-11 22:00:18	3
10	2019-10-11 22:00:28	80
11	2019-10-11 22:00:38	18
12	2019-10-11 22:00:48	267
13	2019-10-11 22:00:58	345
14	2019-10-11 22:01:09	436
15	2019-10-11 22:01:19	698
16	2019-10-11 22:01:29	903
17	2019-10-11 22:01:40	1193
18	2019-10-11 22:01:51	1349
19	2019-10-11 22:02:02	1351
20	2019-10-11 22:02:13	1351

从22:00开始，数据库的活动会话飙升，每秒活动会话飙升至1000以上。故障时间段内的TOP EVENT主要表现在library cache lock和library cache: mutex X等待上。

查看故障期间数据库活动会话情况：

	EVENT	COUNT(*)
1	library cache lock	238555
2	library cache: mutex X	66539
3	cursor: mutex S	9796
4		1019
5	Backup: MML write backup piece	266
6	db file sequential read	171
7	direct path read	133
8	cursor: pin S wait on X	60
9	gc cr disk read	31
10	Backup: MML create a backup piece	28
11	log file sync	18
12	cursor: mutex X	13

49	11-OCT-19 10.00.08.112 PM	...	3794	...	7a732b168c	...	0	0
50	11-OCT-19 10.00.08.112 PM	...	4276	library cache lock	...	g14zrn7wyaxh	4374	5373955 100'mode+namespace
51	11-OCT-19 10.00.08.112 PM	...	4374	cursor: mutex S	...	g14zrn7wyaxh	8168	4294967296 where
52	11-OCT-19 10.00.08.112 PM	...	4498	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
53	11-OCT-19 10.00.08.112 PM	...	4633	library cache: mutex X	...	g14zrn7wyaxh	...	82 where
54	11-OCT-19 10.00.08.112 PM	...	4848	cursor: pin S wait on X	...	g14zrn7wyaxh	...	21474836480 where
55	11-OCT-19 10.00.08.112 PM	...	5199	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
56	11-OCT-19 10.00.08.112 PM	...	5217	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
57	11-OCT-19 10.00.08.112 PM	...	5797	library cache: mutex X	...	g14zrn7wyaxh	...	85 where
58	11-OCT-19 10.00.08.112 PM	...	5906	cursor: mutex S	...	g14zrn7wyaxh	...	4294967296 where
59	11-OCT-19 10.00.08.112 PM	...	6847	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
60	11-OCT-19 10.00.08.112 PM	...	7683	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
61	11-OCT-19 10.00.08.112 PM	...	7919	cursor: mutex S	...	g14zrn7wyaxh	...	4294967296 where
62	11-OCT-19 10.00.08.112 PM	...	8168	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
63	11-OCT-19 10.00.08.112 PM	...	8511	cursor: mutex S	...	g14zrn7wyaxh	...	4294967296 where
64	11-OCT-19 10.00.08.112 PM	...	8741	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
65	11-OCT-19 10.00.08.112 PM	...	8870	cursor: pin S wait on X	...	g14zrn7wyaxh	...	21474836480 where
66	11-OCT-19 10.00.08.112 PM	...	9110	cursor: mutex S	...	g14zrn7wyaxh	...	4294967296 where
67	11-OCT-19 10.00.08.112 PM	...	9325	library cache lock	...	g14zrn7wyaxh	4276	5373954 100'mode+namespace
68	11-OCT-19 10.00.08.112 PM	...	9579	library cache: mutex X	...	g14zrn7wyaxh	...	82 where
69	11-OCT-19 10.00.08.112 PM	...	9692	library cache: mutex X	...	g14zrn7wyaxh	...	68 where
70	11-OCT-19 10.00.08.112 PM	...	9816	kslftc child completion	...	g14zrn7wyaxh	...	0

从10:00:08的ash信息来看，多个library cache lock被4276会话阻塞，4276会话被4374会话“cursor: mutex S”阻塞，同时4374会话被8168“library cache lock”阻塞。从ash分析来看，大量的library cache lock会话的p3值都是5373954和5373955。5373954指的是mode=2，5373955的mode=3，只是持有的方式不同mode=3就是exclusive独占锁。

而4276会话library cache lock的p3值是5373955，对应的namespace HEX:52 —> DEC:82，mode=3。

```
SQL> SELECT indx,kglstdsc FROM x$kgllst where indx=82;
```

```
INDX KGLSTDSC
```

```
82 SQL AREA BUILD
```

@ITPUB博客

SQL AREA BUILD说明library cache lock是在SQL解析上或SQL AREA上的问题。

发生等待是会话都是在执行g14zrn7wyaxh INSERT SQL语句：

```

/** PayOrderMapper.insert */
INSERT INTO TxxxxxxT T
  (T.ID,.....T.Sxxx0)
VALUES
  (SEQ_xxx.nextval,

```

```

:1,:2,:3,:4,:5,:6,:7,:8,:9,:10,:11,:12,:13,:14,:15,:16,:17,:18,:19,:20,:21,:22,:23,:24,:25,:26,:27,:28,:29,:30,:31,:32
,:33,:34,:35,:36,:37,:38,:39,:40,:41,:42,:43,:44,:45,:46,SYSDATE,:47,SYSDATE,:48,:49,:50,:51

```

该SQL中有51个绑定变量，多个绑定变量可能会导致bind variable graduation问题出现，继而导致cursor无法被shared。

INST	BEGIN_SNAP_ID	BEGIN_SNAP_TIME	EXECS	SQL_ID	PLAN	LOADED_VERSIONS	VERSION_COUNT	LOAD	INVALIDATIONS
1	86796	2019-10-11 19:15	140224	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86797	2019-10-11 19:30	138630	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86798	2019-10-11 19:45	128486	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86799	2019-10-11 20:00	159819	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86800	2019-10-11 20:15	100858	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86801	2019-10-11 20:30	112271	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86802	2019-10-11 20:45	117851	g14zxnn7wyaxh	4052461954	2431	35	2	0
1	86803	2019-10-11 21:00	116778	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86804	2019-10-11 21:15	107542	g14zxnn7wyaxh	4052461954	2431	35	2	0
1	86805	2019-10-11 21:30	99224	g14zxnn7wyaxh	4052461954	2431	35	3	0
1	86806	2019-10-11 21:45	518	g14zxnn7wyaxh	4052461954	2515	20	200	21
1	86807	2019-10-11 22:00	4329	g14zxnn7wyaxh	4052461954	3365	89	685	99
1	86808	2019-10-11 22:15	6929	g14zxnn7wyaxh	4052461954	4416	47	477	0
1	86809	2019-10-11 22:30	8264	g14zxnn7wyaxh	4052461954	5373	5	483	0
1	86810	2019-10-11 22:45	39354	g14zxnn7wyaxh	4052461954	5406	38	31	0
1	86811	2019-10-11 23:00	50470	g14zxnn7wyaxh	4052461954	5409	41	5	0
1	86812	2019-10-11 23:15	44148	g14zxnn7wyaxh	4052461954	5410	42	6	0
1	86813	2019-10-11 23:30	20582	g14zxnn7wyaxh	4052461954	5410	42	3	0
1	86814	2019-10-11 23:45	17041	g14zxnn7wyaxh	4052461954	5410	42	3	0
1	86815	2019-10-12 00:00	54091	g14zxnn7wyaxh	4052461954	5410	42	2	0
1	86816	2019-10-12 00:15	37586	g14zxnn7wyaxh	4052461954	5410	42	2	0
1	86817	2019-10-12 00:30	10009	g14zxnn7wyaxh	4052461954	5410	42	2	0
1	86818	2019-10-12 00:45	10196	g14zxnn7wyaxh	4052461954	5411	44	4	0
1	86819	2019-10-12 01:00	40188	g14zxnn7wyaxh	4052461954	5411	44	4	0
1	86820	2019-10-12 01:15	29665	g14zxnn7wyaxh	4052461954	5411	44	2	0
1	86821	2019-10-12 01:30	32731	g14zxnn7wyaxh	4052461954	5411	44	2	0
1	86822	2019-10-12 01:45	27745	g14zxnn7wyaxh	4052461954	5411	44	1	0

从ASH和DBA\_HIST\_SQLSTAT中可以看出21:45分之后SQL频繁load到cursor cache中，其中invalidations有120次，这是从DASH中取的数据，实际数值比采样还要大，另外SQL的LOADED\_VERSION从原来的2431个在短时间内增长到5411个，实际的version count由于11.2.0.3的隐含参数\_cursor\_obsolete\_threshold的关系，version count超过100会重新开始。



这个时候就怀疑是由于SQL的子游标过多引起SQL解析时遍历library cache object handle链表需要很长时间，造成了library cache: mutex x等待。

```
select sql_id,s.version_count,s.loaded_versions from dba_hist_sqlstat s
where LOADED_VERSIONS > 100
and snap_id = (select max(snap_id) from dba_hist_snapshot)
order by 3 desc
;
```

	SQL_ID	VERSION_COUNT	LOADED_VERSIONS
1	g14zxr7wyaxh	48	5415
2	44kjdf6gwsu41	17	5255
3	0nbya8yjkx352	31	4720
4	bygf59cxrw59h	34	3286
5	gfnbw60wryv5w	18	2848
6	27hfx6jhr1tnc	73	2773
7	7q5zktcb1vcxf	69	2368
8	5gj4880d862kz	27	2235
9	qjx8hnp2qr2mg	32	1802

@ITPUB博客

在数据库中可以看出大量loaded\_version超过1000的SQL语句，并且其中有大量游标是过期的。其中SQL\_ID: g14zxr7wyaxh就是此次library cache lock等待最为严重的SQL。

```
select count(*) from v$sql where sql_id='g14zxr7wyaxh' and is_obsolete='Y';
```

	COUNT(*)
1	5367

@ITPUB博客

导致SQL不共享的原因很多，一部分是由于SQL中绑定变量长度不一致导致。

```
select bind_length_upgradeable,count(*)
from v$sql_shared_cursor where sql_id='g14zxr7wyaxh'
group by bind_length_upgradeable
```

	BIND_LENGTH_UPGRADEABLE	COUNT(*)
1	Y	3292
2	N	1948

采集故障期间的AWR，发现当时DB TIME接近2w，平均活动会话达到1200+。

	Snap Id	Snap Time	Sessions	Cursors/Session
Begin Snap:	86807	11-Oct-19 22:00:01	1578	.9
End Snap:	86808	11-Oct-19 22:15:04	1577	1.0
Elapsed:		15.05 (mins)		
DB Time:		19,168.94 (mins)		@ITPUB博客

排在前五的等待事件都属于并发类的等待事件，其中cursor: mutex S等待次数最多。

### Top 5 Timed Foreground Events

Event	Waits	Time(s)	Avg wait (ms)	% DB time	Wait Class
library cache lock	3,930,384	815,626	208	70.92	Concurrency
library cache: mutex X	8,566,339	257,901	30	22.42	Concurrency
cursor: mutex S	9,066,386	73,529	8	6.39	Concurrency
DB CPU		1,199		0.10	
cursor: pin S wait on X	8,270	264	32	0.02	Concurrency

从ASH中分析library cache lock可以得出，多个会话等待library cache lock主要发生在SQL AREA BUILD的mutex持有争用上。Library cache: mutex X 是10.2.0.2之后library cache latch衍生出来等待。

以下是部分等待事件的含义：

#### **cursor: mutex S:**

- We try to get a mutex on Parent cursor or V\$SQLSTAT bucket in shared mode.
- The mutex is "in flux" (someone is in progress of taking it in shared mode) so we have to wait until the holder finishes its shared get.
- Used when:
  - Examining parent cursor, Querying V\$SQLSTATS bucket

#### **library cache: mutex X:**

- Trying to get a mutex on library cache hash bucket in X mode
- The mutex is already held in incompatible mode or is "in flux"

@ITPUB博客

此类等待事件往往都是发生在SQL解析前遍历library cache object handle链表找到shared cursor。

查看AWR中的Mutex Sleep信息发现：Mutex主要有三个函数的sleep是非常高的，kgllkal3 82、kkshGetNextChild[KKSHBKLOC1]、kglUpgradeLock 119。

函数-kgllkal3 82:kgllkal的意思就是kernel generic library cache management library cache lock allocate 82的意思就是SQL AREA BUILD的意思。

函数-kkshGetNextChild [KKSHBKLOC1]:kksh的意思是kernel compile shared objects (cursor) cursor hash table，就是shared cursor的hash链表。持有mutex从library cache 的handle的hash链表上找出可共享的游标。



## Mutex Sleep Summary

- ordered by number of sleeps desc

Mutex Type	Location	Sleeps	Wait Time (ms)
Library Cache	kgllkal3 82	29,580,529	0
hash table	kkshGetNextChild [KKSHBKLOC1]	20,237,948	0
Library Cache	kgIUUpgradeLock 119	10,658,158	0
Library Cache	kgllkd1 85	154,624	0
Library Cache	kgIhdgn2 106	128,465	0
Cursor Pin	kkslce [KKSCHLPIN2]	106,785	0

查看library cache中namespace的命中:

## Library Cache Activity

- "Pct Misses" should be very low

Namespace	Get Requests	Pct Miss	Pin Requests	Pct Miss	Reloads	Invali- dations
ACCOUNT_STATUS	2,614	0.00	0		0	0
BODY	25	0.00	64,158	0.00	0	0
CLUSTER	9	0.00	9	0.00	0	0
DBLINK	2,615	0.00	0		0	0
EDITION	1,565	0.00	2,873	0.00	0	0
HINTSET OBJECT	372	0.00	372	0.00	0	0
INDEX	5,335	0.00	5,335	0.04	2	0
QUEUE	182	0.00	197	0.00	0	0
SCHEMA	1,308	0.00	0		0	0
SQL AREA	48,802	0.47	300,694	6.13	10,740	1,170
SQL AREA BUILD	190,333	0.12	0		0	0
SQL AREA STATS	1,797	62.94	1,797	62.94	0	0
TABLE/PROCEDURE	21,460	0.00	271,263	0.06	117	0
TRANSFORMATION	91	0.00	91	0.00	0	0
TRIGGER	997	0.00	997	0.00	0	0

从AWR中可以看出SQL AREA BUILD被请求次数是最多的，这跟ASH中大量library cache lock是吻合的，SQL AREA中cursro reloads次数也达到10740次。Invalidations达到1170次，说明有很多cursor失效了。

造成library cache lock等一系列严重等待事件的原因是大量的过期游标导致sql解析前花了大量时间去遍历library cache object handle，问题SQL的5415个cursor中有5367个是标记为过期的，查看游标不能被共享的原因：

```
SQL> select count(*) from ssc_reason where reason like '%<reason>Bind mismatch(14)</reason>%';
      3447

SQL> c/14/22
      1* select count(*) from ssc_reason where reason like '%<reason>Bind mismatch(22)</reason>%';
SQL> /
      5257

SQL>
```

@ITPUB博客

造成游标不能被共享的原因中有5257个游标的原因是Bind Mismatch(22)，也就是绑定变量的字符长度发生变化，从32位升级到128位。

```
<reason>Bind mismatch(22)</reason><size>4x4</size><bind_position>40</bind_position><original_oacflg>3</original_oacflg><original_oacmxl>32</original_oacmxl><upgradeable_new_oacmxl>128</upgradeable_new_oacmxl>*
```

```
SQL> select count(*) from ssc_reason where reason like '%<reason>Bind mismatch(14)</reason>%';
      3294

SQL>
```

@ITPUB博客

其中Bind mismatch(14)的也有3294个，这个主要是绑定变量TIMESTAMP类型传值到DATE类型导致的问题。Bind mismatch(14)多发生在第6个绑定变量上，对应表中第7个字段，该字段正好的DATE类型。

```
<ChildNode><ChildNumber>39</ChildNumber><ID>34</ID><reason>Rolling Invalidate Window Exceeded(3)</reason><size>2x4</size><invalidation_window>1570825587</invalidation_window><ksugctm></ChildNode><ChildNode><ChildNumber>39</ChildNumber><ID>40</ID><reason>Bind mismatch(14)</reason><size>4x4</size><bind_position>40</bind_position><original_oacflg>3</original_oacflg><original_oacmxl>32</original_oacmxl><upgradeable_new_oacmxl>128</upgradeable_new_oacmxl></ChildNode>
<ChildNode><ChildNumber>44</ChildNumber><ID>40</ID><reason>Bind mismatch(14)</reason><size>4x4</size><bind_position>40</bind_position><original_oacflg>3</original_oacflg><original_oacmxl>32</original_oacmxl><upgradeable_new_oacmxl>128</upgradeable_new_oacmxl></ChildNode>
153 rows selected.

SQL> 1
      1* select * from ssc_reason where reason like '%<reason>Bind mismatch(14)</reason>%';
SQL>
```

@ITPUB博客

	SQL_ID	DBID	SQL_ID	NAME	POSITION	DUP_POSITION	DATATYPE	DATATYPE_STRING	CHARACTER_SID	PRECISION	SCALE	MAX_LENGTH
1	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
2	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
3	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
4	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
5	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
6	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
7	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
8	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
9	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
10	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
11	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
12	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
13	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
14	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
15	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
16	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11
17	g14zorn7eyanh	3444777292	g14zorn7eyanh	:41	41		160	TIME STAMP			9	11
18	g14zorn7eyanh	3444777292	g14zorn7eyanh	:17	7		160	TIME STAMP				11

综合以上分析,造成大量游标过期的原因有以下两个:

- 1、绑定变量长度导致游标无法共享
- 2、JDBC的bug导致日期类型通过TIMESTAMP传值, 继而导致绑定变量无法共享

相关bug:

Bug 18617175 : JDBC THIN SENDS SCALE VALUE OF 0 OR 9 FOR BINDS CAUSING MANY CHILD CURSORS  
 Bug 12596686 : JDBC THIN APP SENDS SCALE VALUE OF 0 OR 9 FOR BINDS CAUSING MANY CHILD CURSORS  
 Patch 12596686: JDBC THIN APP SENDS SCALE VALUE OF 0 OR 9 FOR BINDS CAUSING MANY CHILD CURSORS

```
This is a PERFECT match for
Bug 12596686 - JDBC thin app sends scale value of 0 or 9 for Timestamp binds
causing many child cursors (Doc ID 12596686.8)

but the customer tried applying the patch
to the client side in production,
using
How to Manually Apply a One-Off Patch to the JDBC Thin Driver (Doc ID
431463.1)

customer was on jdbc thin 11203, so applied the patch to the client side in
production,
same problem.

they then tried applying JDBC 11204 in a test environment and still have the
same problem.
shared cursors showing 0 to 9 and 9 to 0.

so the question is, why all the shared cursors?

thanks
Larry
*** 04/18/14 06:25 am *** (CHG: Sta->16 Sev->1)
*** 04/18/14 06:25 am ***
```

@ITPUB博客

从上面截图的MOS文档来看，JDBC版本升级到11203或11204仍有发生此例绑定变量传值问题。

### 为什么重启应用无法解决？

1、kill session：故障发生后数据库端进行kill session操作，但是因为有连接池，所以连接池会尝试重连数据库，kill 后的重连在连接池上几乎是并发的，因此负载也很高，所以kill session不行；

2、重启应用：重启应用前数据库端的latch竞争一直都有，大量的活动会话并没有释放。如果这个时候重启应用还是会有新的连接进来，这些新进的连接依然会进入到队列中等待，继而加剧争用，因为重启并不会中止数据库上之前的连接，所以重启应用也不行；

3、关闭应用并kill session：应该关闭应用，然后数据库端kill session，再启动应用。

## 问题解决

**建议一：**后期进行历史分区清理的操作（DDL操作同类）时，需提前查询表上SQL的游标是否超过200，如超过这个阈值，应主动使用DBMS\_SHARED\_POOL.PURGE的方式将过期的游标清理出内存，尽可能的减少遍历游标HASH链表时间较长的现象；

查询并清理过期游标的SQL：

```
select q'[exec sys.dbms_shared_pool.purge(']' || address || ',' ||  
      hash_value || q'[','C');]' as flush_sql  
  from v$sql t  
 where t.sql_id = '&sqlid'  
       and t.is_obsolete = 'Y'  
       group by address,hash_value;
```

**建议二：**从应用层面，建议将前述同一个SQLID(g14zxr7wyaxh)的SQL文本，通过在原SQL文本中，加入不同的注释，从而将其变为若干个不同SQLID，但功能相同的SQL。其目的也是业务高峰期时，将访问分散到不同的父游标上。

**其他建议：**

- 1、将单个SQL游标总数加入到监控告警中，前提是v\$sql\_shared\_cursor中的游标总量在阈值内，目前根据测试和经验总结建议阈值设置为200；
- 2、数据库分区维护操作属于DDL操作，影响较大，应选择业务最低峰期进行操作；
- 3、数据库上执行DDL操作时，应实时监控数据库的活动会话等待事件，如果出现mutex或latch等待持续上升，应立即取消DDL操作，并持续监控数据库性能。